



Supply Chain and Logistics, Imperial College Business School
Demand Forecasting for a Fast Food Restaurant
CID: 02005361

Introduction

This assignment is tasked with forecasting how much lettuce should be procured for a fast food restaurant, Subway, in Berkeley California, USA. The data available contains records of over 7,000 orders placed between 15th March 2015 and 15th June 2015, at store 46673. The data tables of interest to us are tables 1 and 2 for order information, tables 4, 5, 6, and 7 for recipe information, tables 8 and 9 for subrecipe information, and tables 10 and 11 for ingredients. Using this dataset our objective is to calculate a forecast of the demand for lettuce over the following two weeks, June 16-29th 2015, using the ARIMA model and the Holt-Winters model.



Data Import & Prep

```
In [1]: # import packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats
from decimal import Decimal
import statsmodels.api as sm
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from sklearn.metrics import mean_squared_error

C:\Users\symond\anaconda3\lib\site-packages\pandas\core\computation\expressions.py:11: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.2' currently installed).
  from pandas.core.computation.check import NUMEXPR_INSTALLED
C:\Users\symond\anaconda3\lib\site-packages\pandas\core\arrays\masked.py:60: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.2' currently installed).
  from pandas.core import (
In [2]: # Import tables
filelist = ('data/1_pos_ordertable.csv', 'data/2_menuitem.csv', 'data/3_store_restaurant.csv', 'data/4_menu_items.csv', 'data/5_recipes.csv', 'data/6_subrecipes.csv', 'data/7_recipetables.csv', 'data/8_recipetables.csv', 'data/9_recipetables.csv', 'data/10_recipetables.csv', 'data/11_recipetables.csv')
for filename in filelist:
    table1 = pd.read_csv(filename, parse_dates=True)
    table2 = pd.read_csv(filelist[0], parse_dates=True)
    table3 = pd.read_csv(filelist[1])
    table4 = pd.read_csv(filelist[2])
    table5 = pd.read_csv(filelist[3])
    table6 = pd.read_csv(filelist[4])
    table7 = pd.read_csv(filelist[5])
    table8 = pd.read_csv(filelist[6])
    table9 = pd.read_csv(filelist[7])
    table10 = pd.read_csv(filelist[8])
    table11 = pd.read_csv(filelist[9])
    table11 = pd.read_csv(filelist[10])

# Filter data for Store ID 46673
orders = table1[table1['StoreNumber'] == 46673]

# Check for duplicates
orders_dups = orders[orders.duplicated(keep=False)]
if orders_dups.empty:
    print("No duplicate records.")

In [3]: ###### COMBINE #####
#####
##### Merge on unique order sale key, MenuItemID, RecipeID
merge2 = pd.merge(table2, orders, on=['HKEY_ORDERSALE', 'StoreNumber'])
merged = pd.merge(table4, merge2, left_on='MenuItemID', right_on='Id')
merged = pd.merge(merged, table3, on='StoreNumber')
# Add date (future attribute)
lettuce_10_table = table10[table10['IngredientName'].str.contains('lettuce', case=False, na=False)]
# Merge combine on IngredientID
lettuce_6_table = pd.merge(merged, lettuce_10_table, on='IngredientID')

#####
##### CALCULATIONS #####
#####
##### Calculate lettuce consumption from recipes
recipe_lettuce_consumption = pd.merge(merged, lettuce_6_table, on='RecipeID')
recipe_lettuce_consumption['TotalQuantity'] = (recipe_lettuce_consumption['Quantity_X'].fillna(0)) * recipe_lettuce_consumption['Factor']
# Calculate lettuce consumption from subrecipes
sub_recipe_lettuce_consumption = pd.merge(merged, table9, on='SubRecipeID')
sub_recipe_ingredients = pd.merge(sub_recipe_table, table9, on='SubRecipeID')
# Join with recipe details for quantities
subrecipe_lettuce_consumption = pd.merge(sub_recipe_ingredients, lettuce_10_table, on='IngredientID')
subrecipe_lettuce_consumption['Factor'] = (subrecipe_lettuce_consumption['Factor'].fillna(0))
subrecipe_lettuce_consumption['Quantity'] = subrecipe_lettuce_consumption['Factor'].apply(Decimal)
# Merge subrecipe with recipe
subrecipe_lettuce_consumption = pd.merge(subrecipe_lettuce_consumption, merged[['RecipeID', 'date_X']], on='RecipeID', how='left')
# Set missing quantity and factor values to 0
subrecipe_lettuce_consumption['Quantity'] = subrecipe_lettuce_consumption['Quantity'].fillna(0)
subrecipe_lettuce_consumption['Factor'] = subrecipe_lettuce_consumption['Factor'].fillna(0)
subrecipe_lettuce_consumption['subrecipe_lettuce_consumption'] = subrecipe_lettuce_consumption.dropna(subset=['date_X'])
# Recalculate consumption
subrecipe_lettuce_consumption['Quantity_X'] = recipe_lettuce_consumption['Quantity_X'].apply(Decimal)
subrecipe_lettuce_consumption['Quantity_Y'] = recipe_lettuce_consumption['Quantity_Y'].apply(Decimal)
subrecipe_lettuce_consumption['TotalQuantity'] = subrecipe_lettuce_consumption['Quantity_X'] + subrecipe_lettuce_consumption['Quantity_Y']
print("Lettuce demand calculations complete.")

#####
##### TABULATE #####
#####
##### Combine recipe and subrecipe consumption
recipe_lettuce_consumption_x = recipe_lettuce_consumption[['date_X', 'TotalQuantity']]
recipe_lettuce_consumption_x['date_X'] = recipe_lettuce_consumption_x.rename(columns={'date_X': 'TransactionDate'})
subrecipe_lettuce_consumption_x = subrecipe_lettuce_consumption[['date_X', 'TotalQuantity']]
subrecipe_lettuce_consumption_x['date_X'] = subrecipe_lettuce_consumption_x.rename(columns={'date_X': 'TransactionDate'})
lettuce_demand_x = pd.concat([recipe_lettuce_consumption_x, subrecipe_lettuce_consumption_x], ignore_index=True)
# Convert date column to datetime
date_format = '%Y-%m-%d'
lettuce_demand['TransactionDate'] = pd.to_datetime(lettuce_demand['TransactionDate'], format=date_format, errors='coerce')
invalid_dates = lettuce_demand['TransactionDate'].isna()
# Drop invalid dates
daily_lettuce_demand = lettuce_demand.groupby('TransactionDate')['TotalQuantity'].sum()
daily_lettuce_demand = daily_lettuce_demand.reset_index()
daily_lettuce_demand.rename(columns={'TotalQuantity': 'Qty (oz)'}, inplace=True)
daily_lettuce_demand_copy = daily_lettuce_demand.copy()
# Add day of week
days_of_week = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
daily_lettuce_demand['Day'] = daily_lettuce_demand['date'].dt.day_name()

#####
##### DAILY SETS #####
#####
# Create datasets for each day to perform weekly seasonality analysis
days_of_week = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
for day in days_of_week:
```

```

    # df = daily_lettuce_demand[df['Day'] == day]
    # select day
    mon_df = days_dataframes(daily_lettuce_demand, 'Monday')
    tue_df = days_dataframes(daily_lettuce_demand, 'Tuesday')
    wed_df = days_dataframes(daily_lettuce_demand, 'Wednesday')
    thu_df = days_dataframes(daily_lettuce_demand, 'Thursday')
    fri_df = days_dataframes(daily_lettuce_demand, 'Friday')
    sat_df = days_dataframes(daily_lettuce_demand, 'Saturday')
    sun_df = days_dataframes(daily_lettuce_demand, 'Sunday')
    print("Median lettuce (oz) on Mondays: " + str(int(mon_df['Qty (oz)'].median())))
    print("Median lettuce (oz) on Tuesdays: " + str(int(tue_df['Qty (oz)'].median())))
    print("Median lettuce (oz) on Wednesdays: " + str(int(wed_df['Qty (oz)'].median())))
    print("Median lettuce (oz) on Thursdays: " + str(int(thu_df['Qty (oz)'].median())))
    print("Median lettuce (oz) on Fridays: " + str(int(fri_df['Qty (oz)'].median())))
    print("Median lettuce (oz) on Saturdays: " + str(int(sat_df['Qty (oz)'].median())))
    print("Median lettuce (oz) on Sundays: " + str(int(sun_df['Qty (oz)'].median())))

#####
##### OUTLINE #####
#####

# check for outliers by day and replace with median
def day_outlier_check(df, new_col, standard_deviations):
    # calculate z-scores
    df[new_col] = (df[new_col] - df[new_col].mean()) / df[new_col].std()
    # calculate median
    median_quantity = df[new_col].median()
    # replace outliers (<score > abs(z)) with the median
    df[new_col] = df2.apply(lambda row: median_quantity if row['z_score'] > 2 or row['z_score'] < -2 else row[new_col], axis=1)
    return df2

mon_df_clean = day_outlier_check(mon_df, 'z_score', 2)
tue_df_clean = day_outlier_check(tue_df, 'z_score', 2)
wed_df_clean = day_outlier_check(wed_df, 'z_score', 2)
thu_df_clean = day_outlier_check(thu_df, 'z_score', 2)
fri_df_clean = day_outlier_check(fri_df, 'z_score', 2)
sat_df_clean = day_outlier_check(sat_df, 'z_score', 2)
sun_df_clean = day_outlier_check(sun_df, 'z_score', 2)

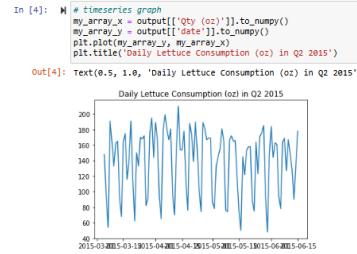
#####
##### OUTPUT #####
#####

# Prepare final data for analysis
daily_dfs = [mon_df_clean, tue_df_clean, wed_df_clean, thu_df_clean, fri_df_clean, sat_df_clean, sun_df_clean]
output = pd.concat(daily_dfs)
output = output.drop(['z_score'], axis=1)
output = output.sort_values(by='date')

<
Lettuce demand calculations complete.
Median lettuce (oz) on Mondays: 174
Median lettuce (oz) on Tuesdays: 154
Median lettuce (oz) on Wednesdays: 157
Median lettuce (oz) on Thursdays: 165
Median lettuce (oz) on Fridays: 94
Median lettuce (oz) on Saturdays: 74
Median lettuce (oz) on Sundays: 164
C:\Users\tsymonds\appdata\local\temp\ipykernel_31968\636268757.py:87: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-or-copy
df[new_col] = (df[new_col] - df[new_col].mean()) / df[new_col].std()
C:\Users\tsymonds\appdata\local\temp\ipykernel_31968\636268757.py:92: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

Timeseries of Lettuce Consumed from March-June 2015 (oz)



ARIMA Analysis

The Autoregressive Integrated Moving Average (ARIMA) forecasting model utilizes autoregressive and moving average components and is a flexible kind of model that can be applied to a broad set of case problems and industries. There are three parameters that we need to pay attention to for this model: p, the count of lagged observations, q, the moving average, and d, the differentiation. We can identify these parameters through the three methods below.

Stationarity

```

In [5]: # create array of consumption
lettuce_data_Q2_2015 = output
lettuce_data_Q2_2015['date'] = pd.to_datetime(lettuce_data_Q2_2015['date'])
lettuce_data_Q2_2015.set_index('date', inplace=True)
lettuce_array = lettuce_data_Q2_2015['Qty (oz)']

In [6]: # check for stationarity
stationary_output = []
result = adfuller(lettuce_array['Qty (oz)'])
stationarity_interpretation = "Stationary" if result[1] < 0.05 else "Non-Stationary"
print("ADF Statistic: " + str(result[0]))
print("p-value: " + str(result[1]))
print("Interpretation: The series is " + stationarity_interpretation)

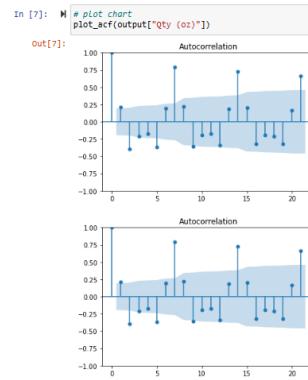
ADF Statistic: -2.242017151518028
p-value: 0.19135022390696556
Interpretation: The series is Non-Stationary.

```

Observations:

We do not have stationarity, which we also could have inferred from the timeseries graph showing a vast discrepancy on a given day. But this check gives us the evidence to estimate our model where d=1. If we did find there to be stationarity, we would have set d=0.

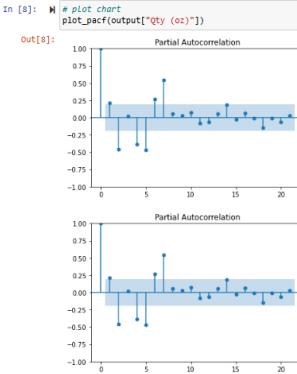
Autocorrelation Function (ACF)



The ACF calculates the correlation between a time series and its lagged values, capturing the degree of similarity between a given time and the time minus one for the ACF to start. This is important for our ARIMA model selection. In general, we want the number of lags before a random walk starts.

say, the current, so yes we performed no var lags would eliminate. In practice, we want some time lags to help reduce a massive spike seen in the chart, and in this case we stop at the first lag, so p=1.

Partial Autocorrelation Function (PACF)



The PACF accounts for the effects of all intermediate lags before calculating the correlation between a time series and its lagged values. We run PACF to get the q parameter for our ARIMA model estimation. Here we also stop at the first lag, so q=1.

ARIMA Model

```
In [9]: # Initial ARIMA Model parameters
p, d, q = 1, 0, 1
# Create the ARIMA model
arima_model = ARIMA(lettuce_array, order=(p, d, q))
# Fit the ARIMA model
arima_result = arima_model.fit()
print("ARIMA Model Fit: " + str(arima_result.bic))
# Forecast the ARIMA model
arima_forecast = arima_result.forecast(steps=14)
# final df
arima_df = pd.DataFrame({'Date':arima_forecast.index, 'ARIMA Forecast (oz)':arima_forecast.values})
print("ARIMA avg forecasted lettuce per day (oz): " + str(arima_df['ARIMA Forecast (oz)'].mean()))

# revised ARIMA Model parameters
p, d, q = 1, 0, 1
# Create the ARIMA model
arima_model = ARIMA(lettuce_array, order=(p, d, q))
# Fit the ARIMA model
arima_result = arima_model.fit()
# Forecast the ARIMA model
arima_forecast = arima_result.forecast(steps=14)
arima_df

C:\Users\lymond\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
C:\Users\lymond\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
C:\Users\lymond\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
C:\Users\lymond\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
BIC Model Fit: 1062.716786720797
ARIMA avg forecasted lettuce per day (oz): 140.59346942206795

C:\Users\lymond\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
C:\Users\lymond\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
C:\Users\lymond\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
C:\Users\lymond\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)

In [10]: print("BIC Model Fit: " + str(arima_result.bic))
print("ARIMA avg forecasted lettuce per day (oz): " + str(arima_df['ARIMA Forecast (oz)'].mean()))
BIC Model Fit: 1062.716786720797
ARIMA avg forecasted lettuce per day (oz): 140.59346942206795
```

Holt-Winters Analysis

Also known as triple exponential smoothing, which refers to the value level, trend, and seasonality behaviors of the time series. The major advantage of the Holt-Winters model over ARIMA is that it can account for seasonality, and does not need too much historical data to be accurate. For our weekly seasonality, we will need to estimate our model with the cadence of seasonality set at 7.

Seasonality - lettuce consumption by day of the week (oz)

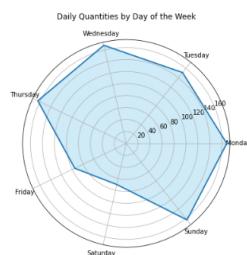
```
In [11]: ##### Day-of-the-week Polar Plot #####
#####
# Calculate daily mean
df_dma = output.copy()
daily_mean_average = df_dma.groupby('Day')[['Qty (oz)']].mean()
daily_mean_list = list(daily_mean_average)

the_daily_list = []
the_daily_means = []
for i, v in daily_mean_average.items():
    the_daily_list.append(i)
    the_daily_means.append(v)
print(the_daily_list)
print(the_daily_means)

# set df
df = pd.DataFrame({'day_of_week': the_daily_list,'quantity': the_daily_means})
# order day of the week
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
df['day_of_week'] = pd.Categorical(df['day_of_week'], categories=day_order, ordered=True)
df = df.sort_values('day_of_week')

# convert day of the week into angles (7 days = 360 degrees)
angles = np.linspace(0, 2 * np.pi, len(df['day_of_week']), endpoint=False).tolist()
# Create a polar plot
quantities = df['quantity'].tolist()
quantities.append(quantities[0]) # Close the loop
quantities.append(angles[0]) # Close the loop

# PLOT the polar chart
fig, ax = plt.subplots(figsize=(6, 6), subplot_kw={'polar': True})
ax.plot(angles, quantities, linewidth=2, linestyle='solid', label='quantity')
ax.fill(angles, quantities, color='skyblue', alpha=0.4)
ax.set_xticks(angles[1::])
ax.set_xticklabels(day_order)
ax.set_title('Daily Quantities by Day of the week', y=1.08)
ax.set_rmax(100)
# Display the chart
plt.show()
```



Observations:

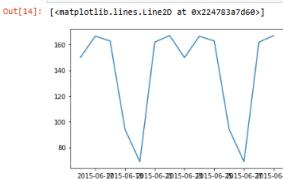
As suggested in our timeseries graph, we anticipated that there would be a difference in average lettuce consumption depending on the day. We can see in the polar chart above that there is a distinct decrease in average daily consumption on Fridays and Saturdays, whereas the demand on the other days is consistent. From this information we can gather that there is weekly seasonality to the demand of lettuce, and this will need to be taken into account in the Holt-Winters model.

Holt-Winters Model

```
In [12]: M # Holt-Winters model
hw_model = ExponentialSmoothing(lettuce_array, seasonal='add', seasonal_periods=7)
hw_result = hw_model.fit()
hw_forecast = hw_result.forecast(steps=14)
hw_forecast_1_year = hw_result.forecast(steps=366)
hw_fcst_df = pd.DataFrame({'Date':hw_forecast.index, 'HW Forecast (oz)':hw_forecast.values})
hw_1yfcst_df = pd.DataFrame({'Date':hw_forecast_1_year.index, 'HW Forecast (oz)':hw_forecast_1_year.values})
print("Holt-Winters avg forecasted lettuce per day (oz): " + str(hw_fcst_df['HW Forecast (oz)'].mean()))
# Fit BIC Model
hw_bic_result = hw_model.fit()
print("BIC Model Fit: " + str(hw_bic_result.bic))
print("Holt-Winters avg forecasted lettuce per day (oz): 138.6948327146055
BIC Model Fit: 631.23353669927
C:\Users\tsa\Downloads\anaconda\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information is provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
```

```
In [13]: M print("Holt-Winters avg forecasted lettuce per day (oz): " + str(hw_fcst_df['HW Forecast (oz)'].mean()))
print("BIC Model Fit: " + str(hw_bic_result.bic))
print("Holt-Winters avg forecasted lettuce per day (oz): 138.6948327146055
BIC Model Fit: 631.23353669927
```

```
In [14]: M # timeseries forecast chart (2 weeks)
hw_array_x = hw_fcst_df[['Date']].to_numpy()
hw_array_y = hw_fcst_df[['HW Forecast (oz)']].to_numpy()
plt.plot(hw_array_x, hw_array_y)
```



Observations:

In forecasting the following two weeks of lettuce demand, the Holt-Winters model is able to project for the Fridays and Saturdays that are going to have much lower consumption than the rest of the weekdays. Based on the clear observed seasonality it makes the Holt-Winters model a good candidate for our final forecast.

Forecast for June 16th-30th 2015

```
In [15]: M # Create 2-week forecast results DataFrame
forecast_dates = pd.date_range(start="2015-06-16", end="2015-06-29")
forecast_results = pd.DataFrame({'Date': forecast_dates, 'ARIMA Forecast (oz)': arima_forecast, 'Holt-Winters Forecast (oz)': hw_forecast})
print(forecast_results)

ARIMA Forecast (oz) Holt-Winters Forecast (oz)
Date
2015-06-16 148.482754 149.040835
2015-06-17 141.767396 166.483812
2015-06-18 140.239810 162.727629
2015-06-19 139.794929 162.727629
2015-06-20 139.812952 68.958894
2015-06-21 139.794962 161.824664
2015-06-22 139.798869 166.925642
2015-06-23 139.798869 166.925642
2015-06-24 139.789726 166.483812
2015-06-25 139.789678 162.727629
2015-06-26 139.789667 94.189352
2015-06-27 139.785664 68.958894
2015-06-28 139.785664 161.824664
2015-06-29 139.785663 166.925642
```

Observations:

Side-by-side we can observe that the ARIMA model provides a constant baseline whereas the Holt-Winters forecast provides a much better prediction with the lower or higher demand figures on a given weekday. The implication of choosing a better model for the forecast lies in saving costs by not buying too much lettuce, in guaranteeing customer satisfaction by not running out of lettuce, and in keeping the ingredient as fresh as possible. Based on these considerations, it is clear that the Holt-Winters model is the optimal forecast to accommodate for the seasonality of the weekdays.

Business Recommendation:

Knowing that Fridays and Saturdays has much less demand for lettuce than any other days of the week, weekly shipments could be delivered on Saturday evenings or Sunday mornings. This ensures the freshest stock for the most consumption during the week, and also leaves the least fresh lettuce for the days where there is much less consumption.

Residual Analysis

Mean Squared Error (MSE) is a measure of the average squared difference of the actual versus predicted values. In other words, this helps us get a better idea of how well our model fits the data. Models with better fits will have lower MSEs.

```
In [16]: M # Get actual lettuce consumed in last two weeks aligned with days of the week occurring during 16th-29th forecast (Tuesday-Tuesday)
twowk_actuals = lettuce_data_Q2_2015.copy()
twowk_actuals = twowk_actuals.tail(14)
twowk_actuals = twowk_actuals.drop('Day', axis=1)
twowk_actuals

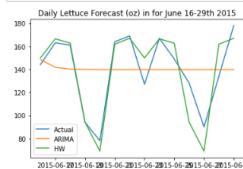
# Create arrays of two-week actuals/forecast for input into residual analysis
twowk_actuals['Day'] = np.array(twowk_actuals)
arima_forecast_array = np.array(arima_forecast)
arima_forecast_array = np.array(arima_forecast)
hw_forecast_array = np.array(hw_forecast)
date_array = np.array(date_array)
date_array = np.array(date_array)

# run for Mean Square Error estimates
arima_mse = mean_squared_error(twowk_actuals_array,arima_forecast_array)
hw_mse = mean_squared_error(twowk_actuals_array,hw_forecast_array)
arima_mse = mean_squared_error(twowk_actuals_array,arima_forecast_array)

print("Mean Squared Error for ARIMA: " + str(arima_mse))
print("Mean Squared Error for Holt-Winters: " + str(hw_mse))
Mean Squared Error for ARIMA: 954.875532440156
Mean Squared Error for Holt-Winters: 238.50086902974085
```

Model Fit

```
In [17]: M # plot actuals, ARIMA forecast, and Holt-Winters forecast
plt.plot(date_array, twowk_actuals_array, label = "Actual")
plt.plot(date_array, arima_forecast_array, label = "ARIMA")
plt.plot(date_array, hw_forecast_array, label = "HW")
plt.legend()
plt.title("Daily Lettuce Forecast (oz) in for June 16-29th 2015")
plt.show()
```



To finalize our forecast, we choose the Holt-Winters model for two reasons. First, the residual analysis suggests it is a far better fitting model than ARIMA, and second, triple exponential smoothing is the optimal method to deal with the heavy weekday seasonability of lettuce consumption. It also gives confidence in our forecast after having visualized the demand forecast on a timeseries chart and observe how the Holt-Winters predictions closely match past consumption.

Final Forecast

```
In [18]: M # 2-week Forecast Table
final_forecast_df = pd.DataFrame(forecast_results['Holt-Winters Forecast (oz)'])
final_forecast_df.reset_index(inplace = True)
final_forecast_df.columns = ['Store', 'California 1 (ID:46673)']
```

```

final_forecast_df
Out[18]:
Store California 1 (ID:46673)
0 2015-06-10 140.840835
1 2015-06-17 166.483012
2 2015-06-18 162.727629
3 2015-06-19 94.103952
4 2015-06-20 68.958094
5 2015-06-21 161.824604
6 2015-06-22 166.025642
7 2015-06-23 140.840835
8 2015-06-24 166.483012
9 2015-06-25 162.727629
10 2015-06-26 94.103952
11 2015-06-27 68.958094
12 2015-06-28 161.824604
13 2015-06-29 166.025642

In [19]: # output to csv
final_forecast_df.to_csv("46673_27_Forecast.csv", index=False)

```

Improving the ARIMA model with revised ACF/PACF lags

In the ACF and PACF charts we noticed that the first lag cutoff is at p=1 and q=1, however, if we look past this we can see than the next noticeable lag is at 7, which matches with the days in the week. When we update our model with p=7 and q=7, we see that the revised ARIMA (rARIMA) forecast is much more accurate, where the bic model fit is much lower and the mean squared error is very close to the Holt-Winters residuals.

```

In [20]: # revised ARIMA Model parameters
p, d, q = 7, 1, 7
# Estimate the ARIMA model
rarma_model = ARIMA(lettuce_array, order=(p, d, q))
# Fit the ARIMA model
rarma_result = rarma_model.fit()
print("BIC Model Fit: " + str(rarma_result.bic))
# Forecast the ARIMA model
rarma_forecast = rarma_result.forecast(steps=14)
# Print of
rarma_df = pd.DataFrame({'Date':rarma_forecast.index, 'rARIMA Forecast (oz)':rarma_forecast.values})
print("ARIMA avg forecasted lettuce per day (oz): " + str(rarma_df['rARIMA Forecast (oz)'].mean()))
print("Mean Squared Error for Revised ARIMA: " + str(rarma_mse))

rarma_df
C:\Users\yusmnd\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
C:\Users\yusmnd\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
C:\Users\yusmnd\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
C:\Users\yusmnd\Anaconda3\lib\site-packages\statsmodels\tsa\state空间\arimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using terms as starting parameters.
warnings.warn('Non-stationary starting autoregressive parameters')
BIC Model Fit: 975.314812899488
rARIMA avg forecasted lettuce per day (oz): 140.18749536544556
Mean Squared Error for Revised ARIMA: 204.59568529447957
C:\Users\yusmnd\Anaconda3\lib\site-packages\statsmodels\tsa\state空间\arimax.py:687: ConvergenceWarning: Maximum likelihood optimization failed to converge. Check mle_retvals
warnings.warn('Maximum likelihood optimization failed to '
"
```

```

Out[20]:
Date rARIMA Forecast (oz)
0 2015-06-16 154.378726
1 2015-06-17 175.510418
2 2015-06-18 165.351343
3 2015-06-19 165.515338
4 2015-06-20 74.287818
5 2015-06-21 145.336834
6 2015-06-22 162.517929
7 2015-06-23 151.110068
8 2015-06-24 175.867900
9 2015-06-25 166.968203
10 2015-06-26 105.988431
11 2015-06-27 75.159173
12 2015-06-28 143.268128
13 2015-06-29 161.492766

```

```

In [21]: print("BIC Model Fit: " + str(rarma_result.bic))
print("rARIMA avg forecasted lettuce per day (oz): " + str(rarma_df['rARIMA Forecast (oz)'].mean()))
print("Mean Squared Error for Revised ARIMA: " + str(rarma_mse))

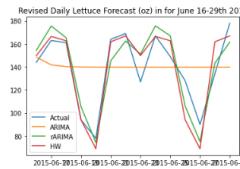
BIC Model Fit: 975.314812899488
rARIMA avg forecasted lettuce per day (oz): 140.18749536544556
Mean Squared Error for Revised ARIMA: 204.59568529447957

```

```

In [22]: # plot actuals, arima forecast, and holt-winters forecast
plt.plot(date_array, book.actuals_array, label = "Actual")
plt.plot(date_array, arima_forecast_array, label = "ARIMA")
plt.plot(date_array, rarma_forecast_array, label = "rARIMA")
plt.plot(date_array, hw_forecast_array, label = "HW")
plt.legend()
plt.title('Revised Daily Lettuce Forecast (oz) in for June 16-29th 2015')
plt.show()

```



```

In [23]: # Revise 2-week Forecast results DataFrame
forecast_dates = pd.date_range(start="2015-06-16", end="2015-06-29")
revised_forecast_results = pd.DataFrame({'Date': forecast_dates, 'ARIMA Forecast (oz)': arima_forecast, 'rARIMA Forecast (oz)': rarma_forecast, 'Holt-Winters Forecast (oz)': hw_forecast})
revised_forecast_results

```

```

Out[23]:
ARIMA Forecast (oz) rARIMA Forecast (oz) Holt-Winters Forecast (oz)
Date
2015-06-16 148.402754 154.378726 140.840835
2015-06-17 141.757366 175.510418 166.483012
2015-06-18 140.239810 165.351343 162.727629
2015-06-19 139.820209 165.515338 94.103952
2015-06-20 139.812052 74.287818 68.958094
2015-06-21 139.794062 145.336834 161.824604
2015-06-22 139.790069 162.517929 166.025642
2015-06-23 139.789933 151.110068 149.840835
2015-06-24 139.789725 175.667900 165.483012
2015-06-25 139.789718 166.062833 162.727629
2015-06-26 139.789567 105.988431 94.103952
2015-06-27 139.789564 75.159173 68.958094
2015-06-28 139.789564 143.268128 161.824604
2015-06-29 139.789563 161.492766 166.025642

```

```

In [24]: # output to csv
revised_forecast_results.to_csv("46673_27_forecast_revised.csv", index=False)

```

```

In [25]: from IPython.core.display import HTML
HTML"""
<style>
div.output_stderr {
background: #e0e0e0;
}
</style>
"""

```